

# CEP Weighted Assessment 3: NoSQL Database

Loh Yu Chen 2020

## Files

- `cli.py`: Main interface for the app
- `library_api.py`: Interface with the MongoDB database
- `config.py`: Credentials for connection to the database
  - `config_template.py`
- `templates.py`: Templates for book and borrower documents
- `insert_books.py`: Script to insert books based on Google Books ID
- `insert_borrowers.py`: Script to insert borrowers from an online API

## Data Model

### Books

#### Fields

- `_id`: ObjectId
- `title`: String
- `authors`: Array (String)
- `isbn`: String
- `page_count`: Int32

#### Indexes

- Text index of `title`, `authors`, `isbn` to allow & speed up search

### Borrowers

#### Fields

- `_id`: ObjectId
- `name`: String
- `username`: String
- `phone`: String

#### Indexes

- Text index of `name`, `username` to allow & speed up search

### Loans

#### Fields

- `_id`: ObjectId
- `book_id`: ObjectId
- `borrower_id`: ObjectId
- `returned`: Boolean

#### Indexes

- Compound index of `book_id`, `returned` as they are used together to check whether a book is currently on loan
- Compound index of `borrower_id`, `returned` as they are used together to check for books that a borrower still has on loan

## Assumptions & Justifications

### Usage pattern

- People will likely use the “checkout” and “return” functions the most as they are the main purpose of a library
  - Hence I allowed multiple books to be loaned/returned at the same time: this allows for more convenience while also reducing the number of operations
- When deleting books/borrowers, some loans may still be active, so all of them are returned automatically
- To show books that exist but are on loan, instead of completely hiding them, I disabled the option to select the book, and displayed the current borrower’s username
  - This makes users aware of the existence of this book, so they can potentially borrow it in the future

### How you are accessing your data

- Most operations are through searches, such as borrowing/deleting/editing books, and creating/editing/deleting borrowers
  - Hence text indexes are very important as they support text search queries on string content
  - Text indexes:
    - Books: title, authors, isbn
    - Borrowers: name, username
- When tracking number of books checked out by a given user, the username is used to identify
  - Hence a username index is used as well

### Which queries are critical in your application

```
db.loans.find_one({"book_id": book_id, "returned": False})
```

- **is\_checked\_out**: This checks the loans collection for active loans on a given book\_id, deeming whether it is currently available to check out, otherwise returning the current borrower who has loaned it
  - This makes use of the compound index, making the query efficient

```
db.loans.find({"borrower_id": borrower_id, "returned": False})
```

- **get\_borrowed\_books:** This gets all active loans of a borrower from the loans collection, for use when creating the “return books” menu, and tracking number of books checked out by a borrower
  - This makes use of the compound index, making the query efficient

`db[collection].find({"$text": {"$search": query}})`

- **search\_function:** This makes use of the text indexes created in the book and borrower collections, to find documents efficiently and accurately

### **Ratios between reads and writes**

- Reads should far exceed writes
  - Other than the create operations, all other operations need to perform at least one query
  - This includes the commonly-used search function, which reads many documents at one go
- Writes are only used upon creating, editing, loaning and returning

### **Amount of information stored and how it grows with time**

- Each loan is stored in a separate document
  - Hence, documents don't grow
    - This ensures there are no problems in older versions of MongoDB
    - Better practice to keep documents a relatively consistent size
  - Loans can be easily extended in the future with more information when needed
    - e.g. Date of borrowing/returning
  - However, this will require an additional query and the server has to search around for documents
    - Hence, indexes are used as much as possible (mentioned earlier) to speed up queries and try to mitigate this impact (though, indexes also come with a cost, especially if the database were to grow larger)

## Reflection

### What are the challenges you faced while doing this project?

- Coming up with a database model: I was conflicted in choosing an appropriate method but I settled for pure linking as I was familiar with it, and it is suitable for a small-scale project like this: Though an extra index has to be created and 2 queries have to be executed, the performance impact is negligible for a small library collection like this project, while making it a lot easier to implement. Hence I thought it was a worthy tradeoff and chose pure linking for this project. Also, because documents don't grow, this model is compatible even with older versions of MongoDB.
- Deciding what to do with missing values and null fields: I was not sure how to handle missing fields after being removed, or fields left empty. In the end, I decided to display both in the app as `<null>`.
- Command line interface: It was difficult to create an easy-to-use interface through just the command line. I ended up using `PyInquirer` to display input questions, allowing arrow navigation and checkbox selections.

### Does knowledge of relational design methods help you in any way, or otherwise?

Yes, especially since I used pure linking - a relational approach which I was already familiar with, thanks to earlier relational database experience and lessons. However, I also had to get used to the document and collection model of NoSQL, compared to the tables and rows in relational databases.

### Do you think this type of application is more suitable to be implemented as a relational DB or non-relational document-based DB like MongoDB? Why?

I think this type of application is more suitable to be implemented as a non-relational database. Firstly, the document model allows new fields and information to be added to books/borrowers at ease, while using a relational table would require schema changes and such. Sometimes, a document may not need the same columns as another too. Also, if this application were to scale, for example to accommodate a larger public library, it would be a lot easier in NoSQL, besides providing superior performance<sup>1</sup>.

### Any other comments?

- Missing fields and null values are both displayed as `<null>`

---

<sup>1</sup> <https://www.mongodb.com/nosql-explained/advantages>

- When creating/editing a new book/borrower, fields left empty will not be set (creating a missing field)
- When deleting a book/borrower, all loans related to it are automatically returned
- The feature “track which user has checked out a book” is embedded in the “checkout” function: when a book that comes up in search is on loan, it will show the username of the borrower
- A limitation of my data model is that a query is required for each book to determine whether it is on loan, which can be very expensive